


This is a copy of a shared ChatGPT conversation
Report conversation


📁 Uploaded a file

check this resume once, then I will give the update and you will make the needed changes after discussions





these all oracle certifications i have currently and i think
these are not expired yet



📁 Uploaded a file

Check out my interview session properly, you will get what they expect from a 3 year experienced employee. I need a lot more practice

Good question. If you have **less than 1 hour**, don't study everything. Focus on the **highest ROI** topics.

1. APIs (5 mins)

Difference between:

GET

- Retrieve data

POST

- Create data

PUT

- Update entire resource

PATCH

- Partial update

DELETE

- Delete resource

REST API

A way for applications to communicate over HTTP.

JSON

```
</> JSON

{
  "name": "Hirak",
  "age": 28
}
```



2. Embeddings (2 mins)

Interview Answer:

Embeddings convert text into numerical vectors so semantically similar texts are placed close together in vector space.

3. Cosine Similarity (2 mins)

Very common.

Cosine similarity measures how similar two vectors are based on the angle between them.

Why?

Higher cosine similarity means more semantic similarity.

4. Chunking (2 mins)

Why chunk?

LLMs cannot process huge documents efficiently. Chunking divides documents into smaller meaningful pieces for better retrieval.

Why 1000-1500 characters?

Balance between context and retrieval accuracy.

5. Why Markdown?

If they ask:

Why convert PDF to Markdown?

Answer:

Markdown preserves headings, lists, tables and document structure better than plain text, improving chunk quality and retrieval.

6. Vector Database

Examples

- Pinecone
- ChromaDB
- FAISS
- Milvus
- Weaviate

Don't say GPT.

7. Prompt Types

Know these names.

- System Prompt

- User Prompt
 - Assistant Prompt
 - Few-shot Prompt
 - Zero-shot Prompt
-

8. Transformer

If asked:

| What is Transformer?

Say:

| Transformer is the deep learning architecture used by modern LLMs. It uses the attention mechanism to understand relationships between words in parallel.

Don't go deeper unless asked.

9. GPT

GPT = **Generative Pre-trained Transformer**

GPT is

✓ Decoder-only model

Why?

| Because it predicts the next token.

10. Encoder vs Decoder

Encoder

- Understands input

Decoder

- Generates output

GPT

- Decoder

BERT

- Encoder

T5

- Encoder + Decoder
-

11. Fine-tuning vs RAG

Fine-tuning

- Changes model weights
- Teaches new behaviour

RAG

- Doesn't change model
 - Retrieves external knowledge
-

12. Agentic AI

Chatbot

Answers questions

Agent

Can

- Think
 - Plan
 - Call APIs
 - Use tools
 - Execute workflows
-

13. Why AI instead of Python Automation?

This was asked in your interview.

Good answer:

"If the workflow is rule-based and deterministic, Python automation is sufficient. We use GenAI when the task requires understanding natural language, reasoning over unstructured documents, summarization, or conversational interaction."

This answer is much stronger than saying "AI was decided by the team."

14. Testing

You mentioned testing.

Remember these:

- Functional Testing
 - Prompt Testing
 - Retrieval Validation
 - Hallucination Testing
 - User Acceptance Testing (UAT)
-

15. Python Functions

Know:

Python



Run

```
def add(a,b):  
    return a+b
```

Difference

`return`

Returns value

`print`

Only displays value

★★★ Most Likely TCS Questions

Based on your previous interviews, I would bet on these:

1. Tell me about yourself.
2. Explain your KT Assistant.
3. Explain RAG.
4. Why embeddings?
5. Why vector database?

6. How did you reduce hallucinations?
 7. What was your role?
 8. Python coding (largest, frequency, average, string reverse, dictionary)
 9. Mutable vs Immutable.
 10. `append()` vs `extend()`.
 11. `sort()` vs `sorted()`.
 12. Time complexity.
 13. Agent vs Chatbot.
 14. Fine-tuning vs RAG.
 15. REST API basics.
-



Last 20 Minutes

- ✗ Don't open YouTube.
- ✗ Don't read articles.
- ✗ Don't learn new frameworks.

Only revise:

- Your project
 - Python basics
 - RAG flow
 - GenAI fundamentals
-

One last reminder

If they ask something you don't know, don't panic. Say:

"I haven't worked with that directly, but based on my understanding..."

That sounds much more professional than "I don't know."

You're in a much better position today than you were before the Infosys interview. Trust your experience, keep your answers structured, and take a second to think before you speak.

Based on everything you've shared, I don't think it's possible to give a meaningful numerical probability like "70%" or "80%". Interview outcomes depend on factors we can't observe, such as the interviewer's scorecard, other candidates, the team's hiring needs, and budget.

What I can do is assess your interview performance.

My assessment

Technical fit: Good (around 7.5–8/10)

Reasons:

-  You had real project experience instead of only theoretical knowledge.
-  You answered confidently about prompt engineering, testing, hallucination reduction, and UAT.
-  You communicated well and stayed composed.
-  The interviewer spent a long time giving you constructive feedback rather than ending the interview quickly.

Areas that likely reduced your chances

-  LLM fundamentals weren't deep enough.
-  You mixed concepts like LLM, embeddings, and RAG during explanations.
-  Some answers lacked the "why" behind the design choices.
-  You couldn't discuss alternatives to your RAG implementation in depth.

Positive signs

The interviewer said things like:

- "You have good understanding."
- "You have worked on it."

- "Explore more."
- "Don't think you need to know everything."

That doesn't sound like someone dismissing your profile. It sounds like someone who sees a candidate with relevant experience who needs stronger fundamentals.

My expectation

If I had to summarize your chances:

- **There is a realistic chance of selection.**
- **There is also a realistic chance of rejection if another candidate demonstrated deeper LLM fundamentals.**

I would not walk away thinking, "I definitely failed." I would also not assume, "I'm definitely selected."

The bigger picture

After reviewing your:

- LTIMindtree interview,
- Infosys interview, and
- TCS interview,

I see a clear trend:

- **LTIMindtree:** Python fundamentals exposed.
- **Infosys:** RAG depth exposed.
- **TCS:** Much better communication, but LLM depth exposed.

That's actually encouraging, because each interview has been stronger than the previous one, and the gaps have become more specific.

So my advice is: **don't wait for the TCS result before continuing your preparation.** Use the interviewer's feedback as a roadmap. If you strengthen Python fundamentals and deepen your understanding of LLMs and RAG over the next couple of weeks, you'll be in a much stronger position for the next GenAI interview, whether it's TCS or another company.

Let's continue on the preparation then

Perfect. We now have a clear goal.

We're no longer preparing to "clear interviews." We're preparing to become someone who can confidently answer **deep technical follow-up questions.**

From now on, I'll behave exactly like a senior interviewer.

Topic 1: LLM Fundamentals (Most Important)

Question 1

You said in the interview:

"LLM forms the answer."

Now I ask:

What exactly does an LLM do?

Don't explain RAG.

Don't explain embeddings.

Only explain **what an LLM is** and **what its job is**.

Answer as if you're sitting in front of the interviewer.

Take your time. I won't interrupt.

Llm connects with the user with a simple language, forms the answer in a natural language processing method and gives it back so that user can easily understand it.

Good attempt. You're improving, but this answer would still lose marks.

Score: 5/10

What you said

"LLM connects with the user with a simple language..."

That's more about the **user experience** than what an LLM actually is.

Interview Answer (8-9/10)

"An LLM (Large Language Model) is a deep learning model, based on the Transformer architecture, that is trained on a large amount of text data. Its primary job is to understand the input context and generate human-like text by predicting the next most probable token. In a RAG application, the LLM uses the retrieved context to generate a relevant and natural language response."

Notice how this answer:

- Defines **what it is**
 - Explains **how it works (next-token prediction)**
 - Explains **its role in your project**
-

If the interviewer asks:

Why do we need an LLM? Why can't we just return the retrieved chunks?

What would you answer?

Think about your KT Assistant. Don't worry if it's not perfect. This is exactly the kind of follow-up that TCS and Infosys interviewers ask.